

Deep Generative Models

Affective Computing

Taher Ahmadi

Summer 2020

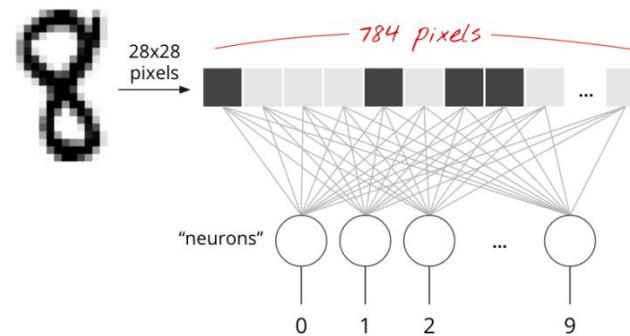
What's a NN?

- Let's start with an example:
- Handwritten digits in the [MNIST](#) dataset are 28x28 pixel grayscale images.

0 0 0 0 0 0 0 0 0
1 1 1 1 1 1 1 1 1
2 2 2 2 2 2 2 2 2
3 3 3 3 3 3 3 3 3
4 4 4 4 4 4 4 4 4
5 5 5 5 5 5 5 5 5
6 6 6 6 6 6 6 6 6
7 7 7 7 7 7 7 7 7
8 8 8 8 8 8 8 8 8
9 9 9 9 9 9 9 9 9

One-Layer Network For Classifying MNIST

- Let's make a **one-layer neural network** for **classifying digits**.
- Each **neuron** in a neural network:
 - Does a **weighted sum** of all of its inputs
 - Adds a **bias**
 - Feeds the result through some **non-linear activation** function, e.g., **softmax**.

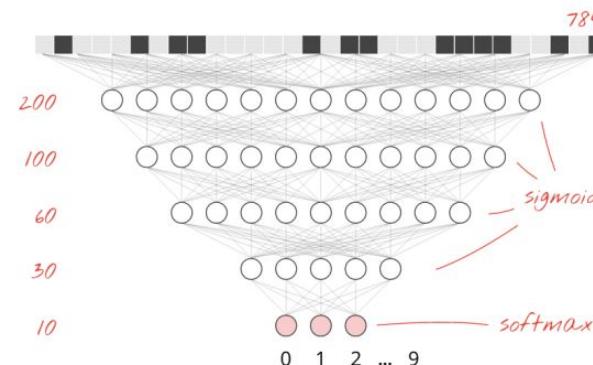


Vanilla Deep Neural Network

- Add more layers to improve the accuracy.
- On intermediate layers we will use the the **sigmoid** activation function.
- We keep **softmax** as the activation function on the last layer.

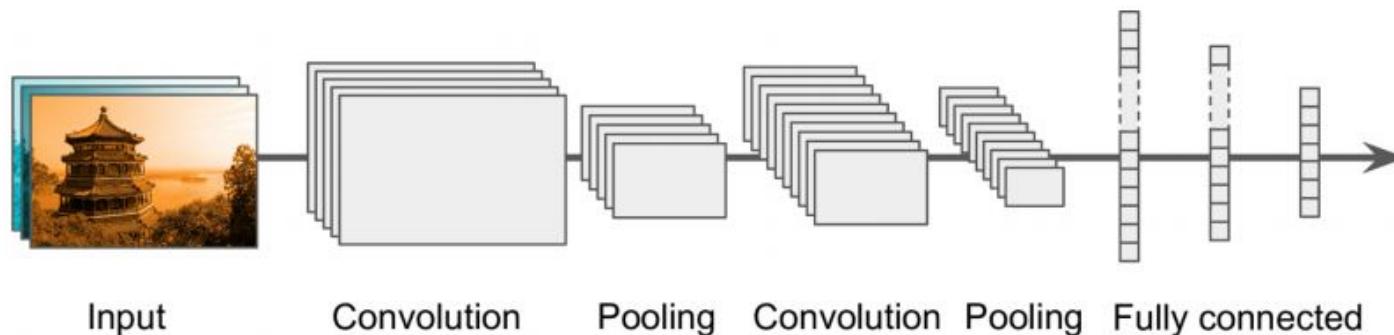


[<https://github.com/GoogleCloudPlatform/tensorflow-without-a-phd>]

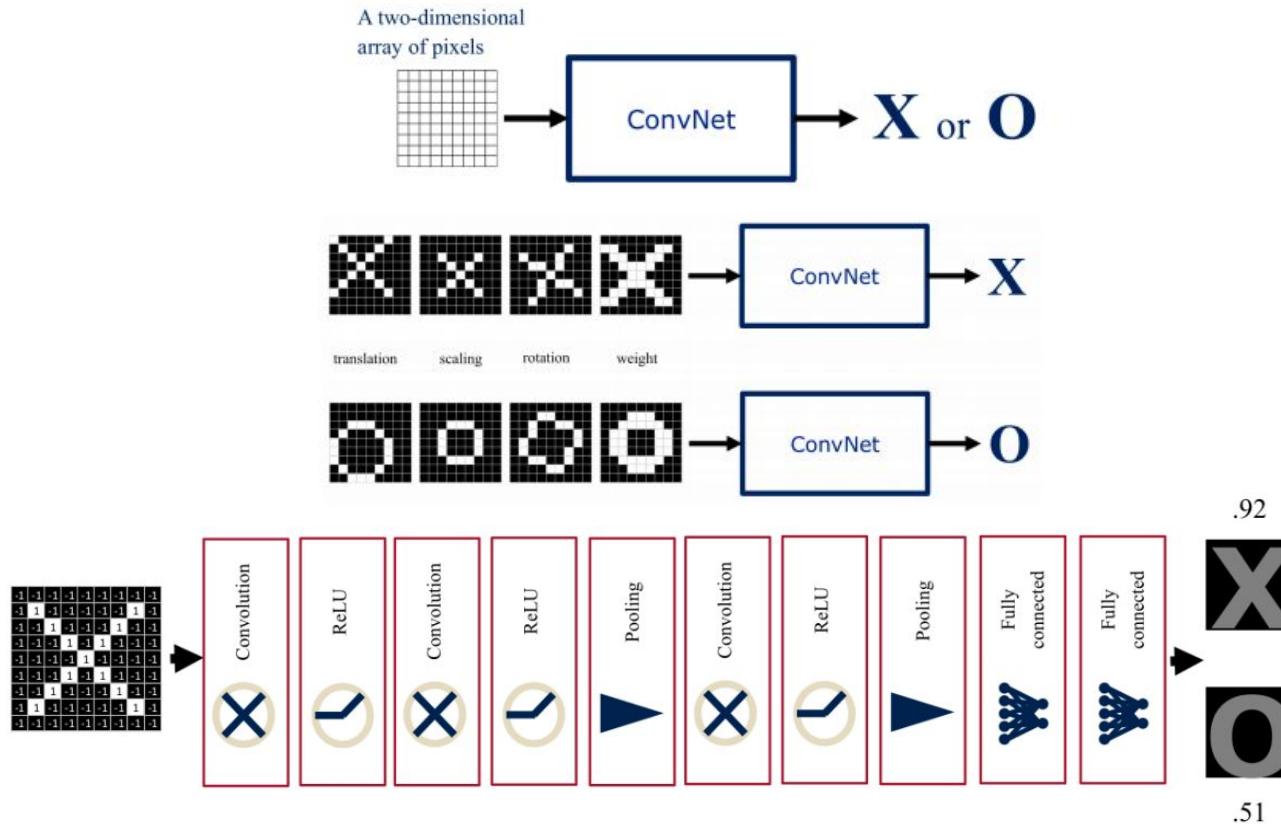


Convolutional Neural Networks

- **Convolutional layers:** apply a specified number of **convolution filters** to the image.
- **Pooling layers:** **downsample the image** data extracted by the convolutional layers to **reduce the dimensionality** of the feature map in order to decrease processing time.
- **Dense layers:** a **fully connected layer** that performs **classification** on the features extracted by the convolutional layers and downsampled by the pooling layers.



CNN for X/O



Neural Networks in Keras

Keras

- Keras is a high-level API to build and train deep learning models.
- To get started, import `tf.keras` to your python code.

```
import tensorflow as tf
from tensorflow.keras import layers
```

Keras Layers

- In Keras, you assemble layers `tf.keras.layers` to build models.
- A model is (usually) a `graph of layers`.
- There are many types of layers, e.g., `Dense`, `Conv2D`, `RNN`, ...

Keras Layers

- Common constructor parameters:
 - activation: the activation function for the layer.
 - kernel initializer and bias initializer: the initialization schemes of the layer's weights.
 - kernel regularizer and bias regularizer: the regularization schemes of the layer's weights, e.g., L1 or L2.

```
layers.Dense(64, activation=tf.sigmoid, kernel_regularizer=tf.keras.regularizers.l1(0.01),  
bias_initializer=tf.keras.initializers.constant(2.0))
```

Keras Models

- There are **two ways** to build Keras **models**: **sequential** and **functional**.
- The **sequential API** allows you to create models **layer-by-layer**.
- The **Functional API** allows you to create models that have a lot **more flexibility**.
 - You can define models where layers connect to more than just their previous and next layers.

Keras Models - Sequential Models

- You can use `tf.keras.Sequential` to build a **sequential model**.

```
from tensorflow.keras import layers

model = tf.keras.Sequential()

model.add(layers.Dense(64, activation="relu"))
model.add(layers.Dense(64, activation="relu"))
model.add(layers.Dense(10, activation="softmax"))
```

Keras Models - Functional Models

- You can use `tf.keras.Model` to build a **functional model**.

```
from tensorflow.keras import layers

inputs = tf.keras.Input(shape=(32,32))
x = layers.Dense(64, activation="relu")(inputs)
x = layers.Dense(64, activation="relu")(x)
predictions = layers.Dense(10, activation="softmax")(x)

model = tf.keras.Model(inputs=inputs, outputs=predictions)
```

Training Keras Models

- Call the `compile` method to configure the learning process.
- `tf.keras.Model.compile` takes three important arguments.
 - `optimizer`: specifies the training procedure.
 - `loss`: the cost function to minimize during optimization, e.g., mean squared error (`mse`), `categorical_crossentropy`, and `binary_crossentropy`.
 - `metrics`: used to monitor training.
- Call the `fit()` method to fit the model the training data.

```
model.compile(optimizer=tf.train.GradientDescentOptimizer(0.001), loss="mse", metrics=["mae"])

model.fit(training_data, training_labels, epochs=10, batch_size=32)
```

Evaluate and Predict

- `tf.keras.Model.evaluate`: evaluate the cost and metrics for the data provided.
- `tf.keras.Model.predict`: predict the output of the last layer for the data provided.

```
model.evaluate(test_data, test_labels, batch_size=32)

model.predict(test_data, batch_size=32)
```

Feedforward Network in Keras

- Building and Training a two layer **sequential model** in Keras.

```
n_neurons_h = 4
n_neurons_out = 3
n_epochs = 100
learning_rate = 0.1

model = tf.keras.Sequential()

model.add(layers.Dense(n_neurons_h, activation="sigmoid"))
model.add(layers.Dense(n_neurons_out, activation="sigmoid"))

model.compile(optimizer=tf.train.GradientDescentOptimizer(learning_rate=.001),
              loss="binary_crossentropy", metrics=["accuracy"])
model.fit(training_X, training_y, epochs=n_epochs)
```

RNNs

An example...



- **Language modeling** is the task of predicting what word comes next

n-gram Language Models

- The students have to ...
- How to learn a Language Model?
- Learn a n-gram Language Model!
- A n-gram is a chunk of n consecutive words.
 - Unigrams: "the", "students", "have", "to"
 - Bigrams: "the students", "students have", "have to"
 - Trigrams: "the students have", "students have to"
 - 4-grams: "the students have to"
 - ...
- Collect statistics about how frequent different n-grams are, and use these to predict next word.

$p(w_j | \text{students have to}) = \text{students have to } w_j / \text{students have to}$

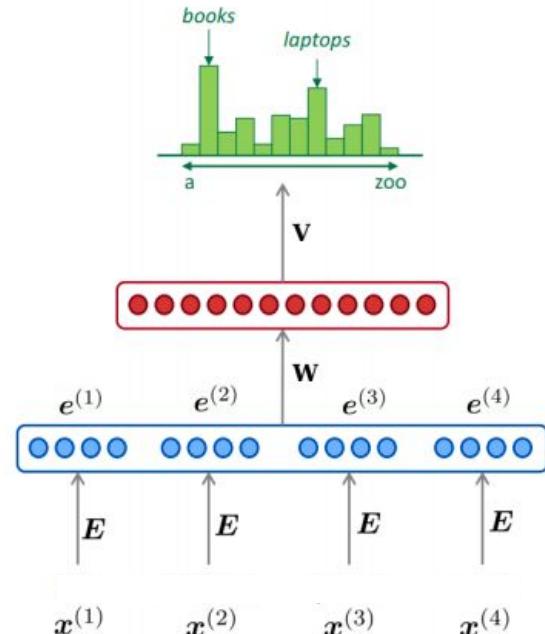
Problems with n-gram Language Models

- What if "**students have to** w_j " never occurred in data?
 - Then w_j has probability 0!
- What if "**students have to**" never occurred in data?
 - Then we can't calculate probability for any w_j !
- Increasing n makes **sparsity** problems worse.
 - Typically we can't have n bigger than 5.
- Increasing n makes model size **huge**.

Can we use a Neural Model?

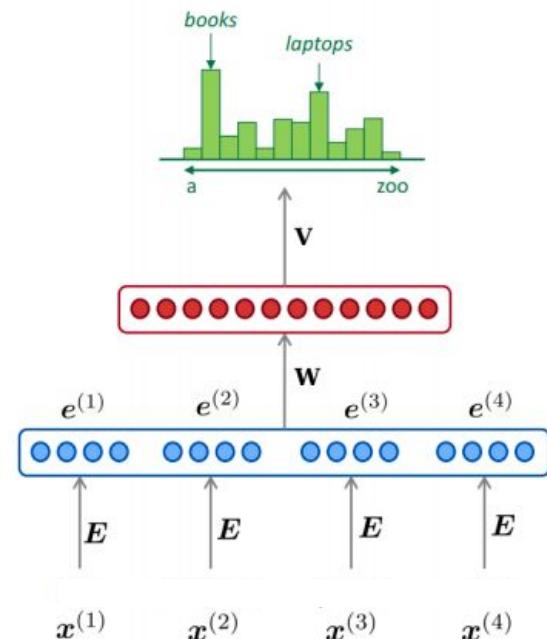
Recall the [Language Modeling](#) task:

- **Input:** sequence of words $x^{(1)}, x^{(2)}, \dots, x^{(t)}$
- **Output:** probability dist of the next word
 $p(x^{(t+1)} = w_j | x^{(t)}, \dots, x^{(1)})$
- MLP model:
 - Input: words $x^{(1)}, x^{(2)}, x^{(3)}, x^{(4)}$
 - Input layer: one-hot vectors $e^{(1)}, e^{(2)}, e^{(3)}, e^{(4)}$
 - Hidden layer: $\mathbf{h} = f(\mathbf{w}^T \mathbf{e})$, f is an activation function.
 - Output: $\hat{\mathbf{y}} = \text{softmax}(\mathbf{v}^T \mathbf{h})$



Can we use a Neural Model?

- **Improvements over n-gram LM:**
 - No sparsity problem
 - Model size is $O(n)$ not $O(\exp(n))$
- **Remaining problems:**
 - It is **fixed** 4 in our example, which is small
 - We need a neural architecture that can process **any length input**

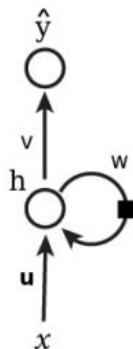


Recurrent Neural Networks (RNN's)

- The idea behind **Recurrent neural networks (RNN)** is to make use of **sequential data**.
 - Until here, we assume that **all inputs (and outputs)** are **independent** of each other.
 - It is a **bad idea** for many tasks, e.g., **predicting the next word in a sentence** (it's better to know which words came before it).
- They can analyze **time series data** and **predict the future**.
- They can work on **sequences of arbitrary lengths**, rather than on **fixed-sized inputs**.

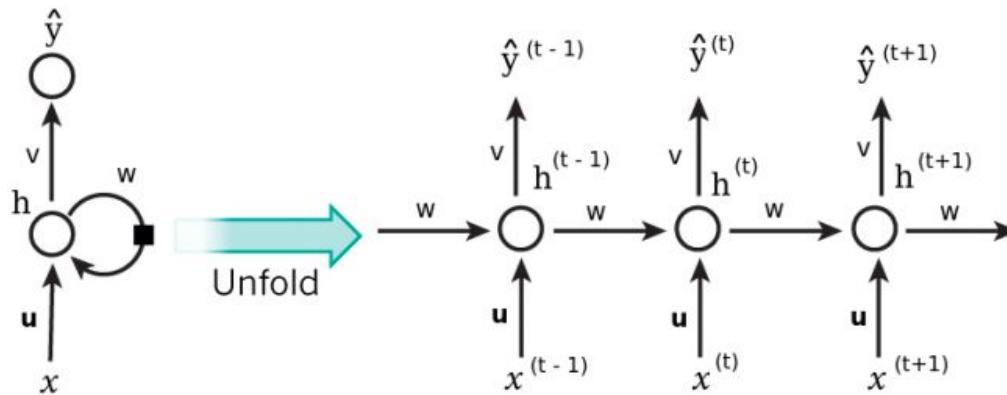
Recurrent Neural Networks

- Neurons in an **RNN** have **connections pointing backward**.
- RNNs have **memory**, which captures **information about what has been calculated so far**.



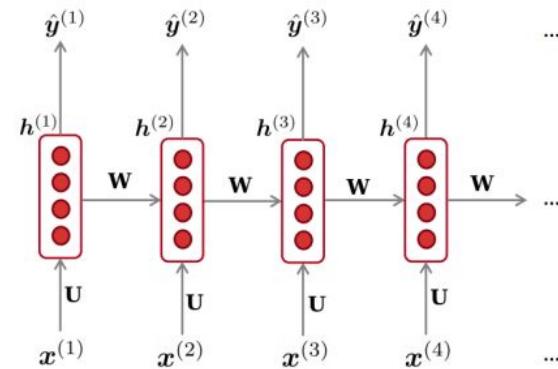
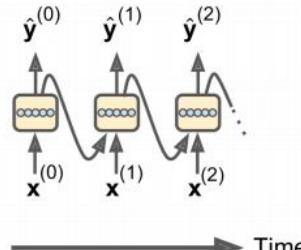
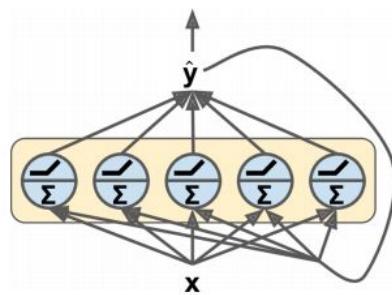
Recurrent Neural Networks

- **Unfolding the network:** represent a network against the time axis.
 - We write out the network for the complete sequence.
- For example, if the sequence we care about is a sentence of three words, the network would be unfolded into a 3-layer neural network.
 - One layer for each word.



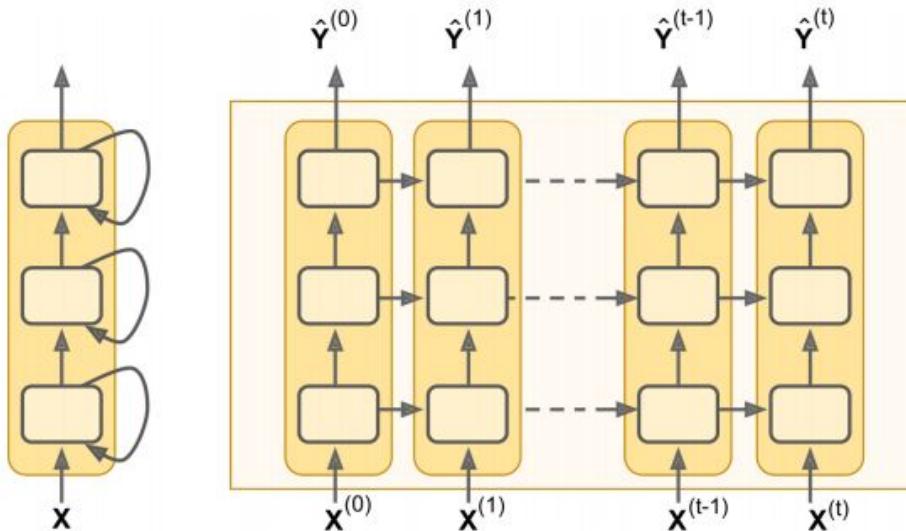
Layers of Recurrent Neurons

- At each time step t , every neuron of a layer receives both the **input vector** $x^{(t)}$ and the **output vector** from the previous time step $h^{(t-1)}$.



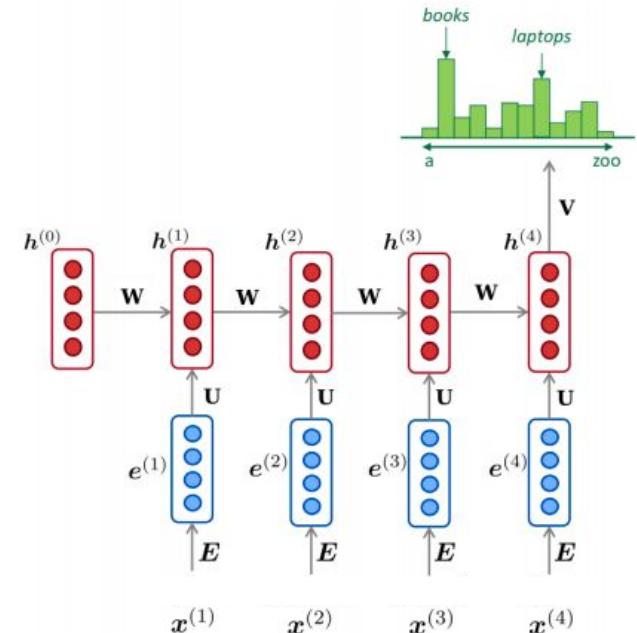
Deep RNN

- Stacking **multiple layers** of cells gives you a **deep RNN**.



Let's Back to Language Model Example

- The equations for the RNN:
 - $h^{(t)} = \tanh(u^t e^{(t)} + wh^{(t-1)})$
 - $\hat{y}^{(t)} = \text{softmax}(vh^{(t)})$
- The output $\hat{y}^{(t)}$ is a vector of **vocabulary size** elements.
- Each element of $\hat{y}^{(t)}$ represents the **probability** of that word being the next word in the sentence.

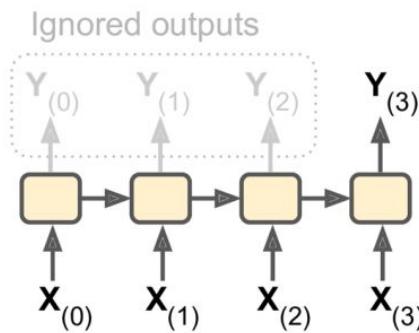


RNN Design Patterns

- Sequence-to-vector
- Vector-to-sequence
- Sequence-to-sequence
- Encoder-decoder

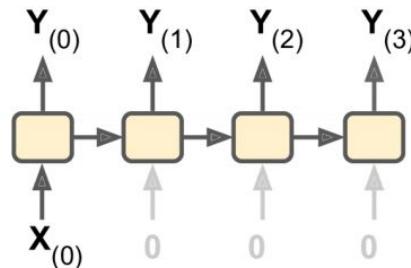
RNN Design Patterns

- **Sequence-to-vector** network: takes a **sequence of inputs**, and ignore all outputs except for **the last one**.
- E.g., you could feed the network a **sequence of words** corresponding to a movie review, and the network would output a **sentiment score**.



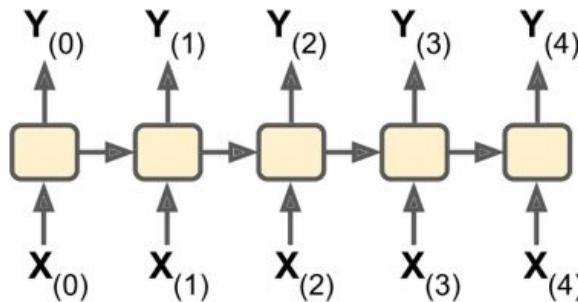
RNN Design Patterns

- **Vector-to-sequence** network: takes a **single input** at the first time step, and let it **output a sequence**.
- E.g., the input could be an **image**, and the output could be a **caption for that image**.



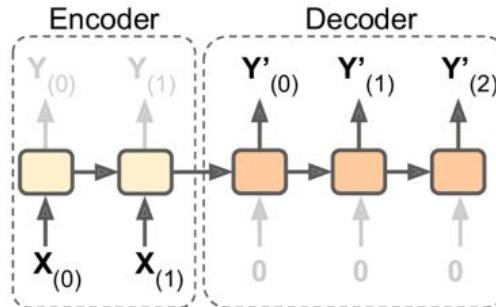
RNN Design Patterns

- **Sequence-to-sequence** network: takes **a sequence of inputs** and produce **a sequence of outputs**.
- Useful for **predicting time series such as stock prices**: you feed it the prices over the last N days, and it must output the prices shifted by one day into the future.
- Here, both input sequences and output sequences have the **same length**.



RNN Design Patterns

- **Encoder-decoder** network: a sequence-to-vector network (**encoder**), followed by a vector-to-sequence network (**decoder**).
- E.g., translating a sentence from one language to another.
- You would feed the network a sentence in one language, the encoder would convert this sentence into a single vector representation, and then the decoder would decode this vector into a sentence in another language.



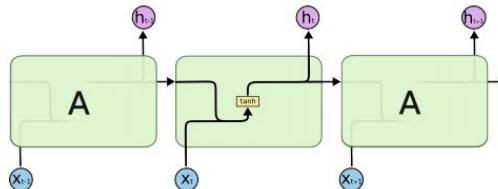
LSTMs

RNN Problems

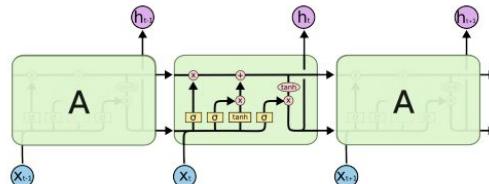
- Sometimes we only need to look at **recent information** to perform the present task.
 - E.g., **predicting** the next word based on the previous ones.
- In such cases, where the **gap between the relevant information and the place that it's needed** is **small**, RNNs can learn to use the past information.
- But, as that **gap grows**, RNNs become **unable to learn to connect the information**.
- RNNs may suffer from the **vanishing/exploding gradients problem**.
- To solve these problem, **long short-term memory (LSTM)** have been introduced.
- In **LSTM**, the network can learn **what to store** and **what to throw away**.

Looking inside LSTM

- The **LSTM** cell looks exactly like a **basic RNN** cell.
- The repeating module in a **standard RNN** contains a **single layer**.



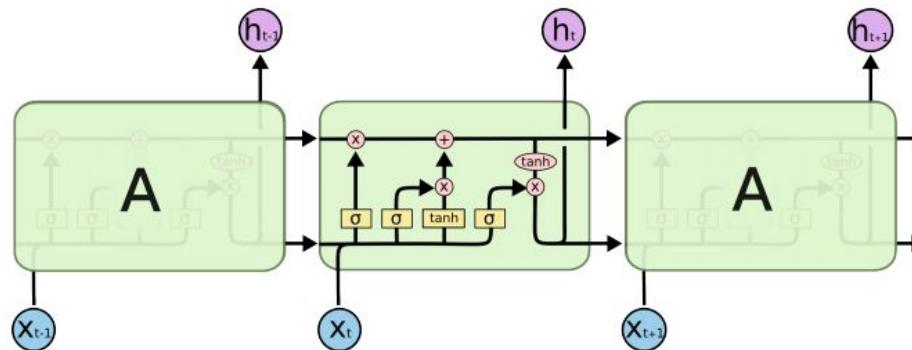
- The repeating module in an **LSTM** contains **four interacting layers**.



LSTM cell

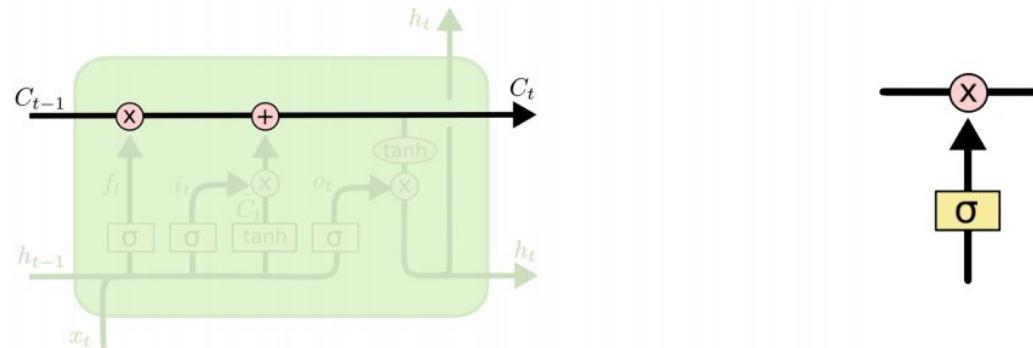
In LSTM **state** is split in two vectors:

1. $h^{(t)}$ (h stands for **hidden**): the **short-term** state
2. $c^{(t)}$ (c stands for **cell**): the **long-term** state



LSTM cell

- The **cell state** (long-term state), the horizontal line on the top of the diagram.
- The LSTM can **remove/add** information to the **cell state**, regulated by three **gates**:
 - **Forget gate, input gate and output gate**



Autoencoders

An Example

- ▶ Which of them is easier to memorize?
- ▶ Seq1: 40, 27, 25, 36, 81, 57, 10, 73, 19, 68
- ▶ Seq2: 50, 25, 76, 38, 19, 58, 29, 88, 44, 22, 11, 34, 17, 52, 26, 13, 40, 20

Seq1 : 40, 27, 25, 36, 81, 57, 10, 73, 19, 68

Seq2 : 50, 25, 76, 38, 19, 58, 29, 88, 44, 22, 11, 34, 17, 52, 26, 13, 40, 20

- Seq1 is shorter, so it should be easier.
- But, Seq2 follows two simple rules:
 - Even numbers are followed by their half.
 - Odd numbers are followed by their triple plus one.
- You don't need pattern if you could quickly and easily memorize very long sequences
- But, it is hard to memorize long sequences that makes it useful to recognize patterns.

Memory

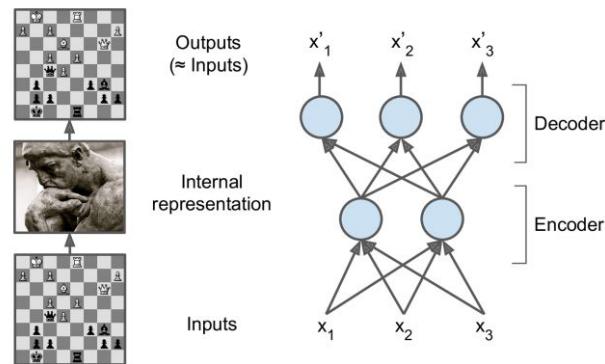
- 1970, W. Chase and H. Simon
- They observed that **expert chess players** were able to **memorize** the positions of all the pieces in a game by looking at the board for just **5 seconds**.



- This was only the case when the **pieces were placed** in **realistic positions**, not when the pieces were placed **randomly**.
- Chess experts **don't have a much better memory** than you and I.
- They just see chess **patterns more easily** due to their **experience** with the game. **Patterns** helps them store information efficiently.

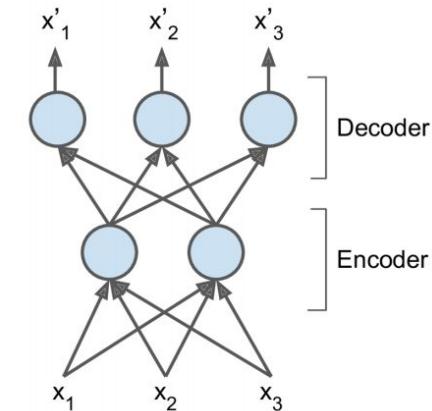
Autoencoders

- Just like the chess players in this memory experiment.
- An **autoencoder** looks at the inputs, **converts** them to an **efficient internal representation**, and then **spits out** something that **looks very close** to the inputs.



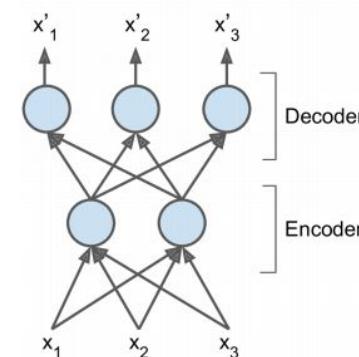
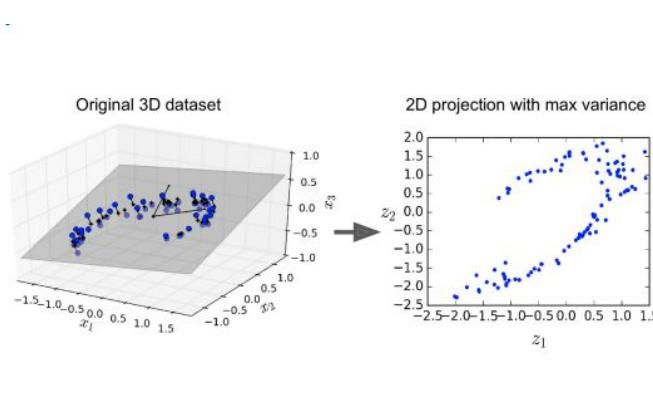
Autoencoders

- An **autoencoder** is always composed of two parts.
- An **encoder (recognition network)**, $h = f(x)$
Converts the **inputs** to an internal representation.
- A **decoder (generative network)**, $r = g(h)$
Converts the **internal representation** to the **outputs**.
- If an autoencoder learns to set $g(f(x)) = x$ everywhere, it is **not especially useful**, **why?**
 - Autoencoders are designed to be **unable to learn to copy** perfectly.
 - The models are forced to prioritize **which aspects of the input should be copied**, they often learn **useful properties of the data**.



Autoencoders

- **Autoencoders** are neural networks capable of learning **efficient representations of the input data** (called **codings**) without any supervision.
- **Dimension reduction:** these **codings** typically have a much lower **dimensionality** than the **input data**.

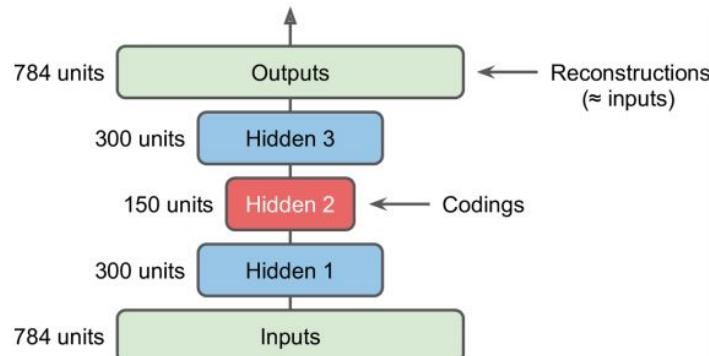


Two Types of Autoencoders

- Stacked autoencoders
- Variational autoencoders

Stacked autoencoders

- **Stacked autoencoder**: autoencoders with **multiple hidden layers**.
- Adding **more layers** helps the autoencoder learn more **complex codings**.
- The architecture is typically **symmetrical** with regards to the **central hidden layer**.

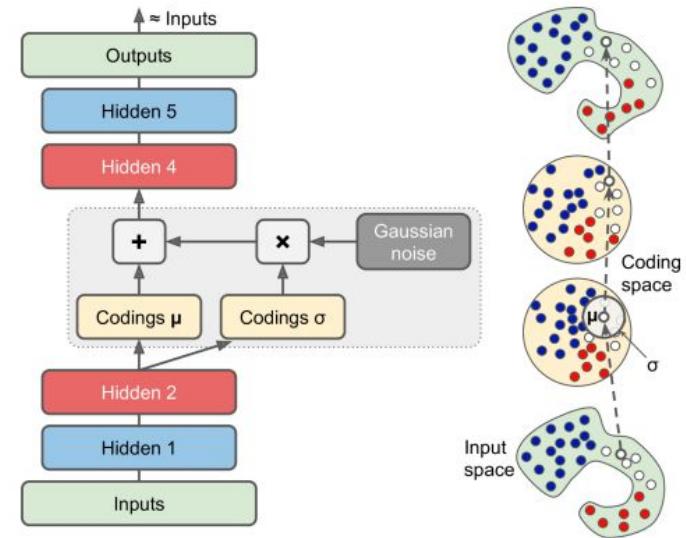


Variational Autoencoders

- Variational autoencoders are probabilistic autoencoders.
- Their outputs are partly determined by chance, even after training.
 - As opposed to denoising autoencoders, which use randomness only during training.
- They are generative autoencoders, meaning that they can generate new instances that look like they were sampled from the training set.

Variational Autoencoders

- Instead of directly producing a coding for a given input, the **encoder** produces a **mean coding μ** and a **standard deviation σ** .
- The **actual coding** is then **sampled randomly** from a **Gaussian distribution** with **mean μ** and **standard deviation σ** .
- After that the **decoder** just **decodes the sampled coding normally**.

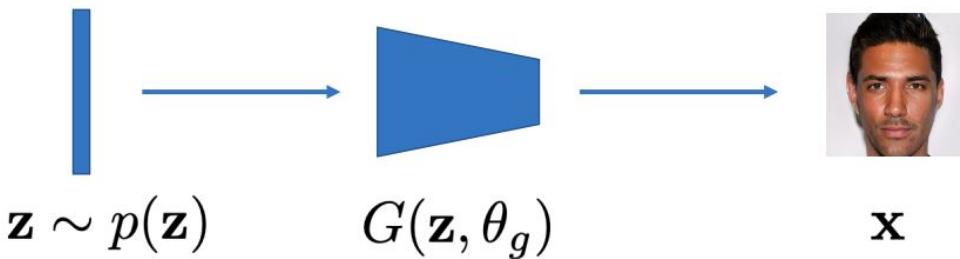


Variational Autoencoders

- The **cost function** is composed of **two parts**.
 1. the usual **reconstruction loss**.
 - Pushes the autoencoder to **reproduce its inputs**.
 - Using **cross-entropy**
 2. the **latent loss**
 - Pushes the autoencoder to have **codings** that look as though they were sampled from a simple **Gaussian distribution**.
 - Using the **KL divergence** between the **target distribution** (the Gaussian distribution) and the **actual distribution** of the codings.
 - KL divergence measures the **divergence between the two probabilities**.

Generative adversarial networks

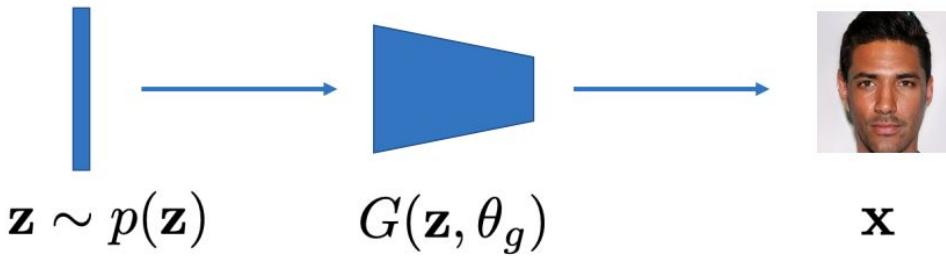
GANs



Sample a random code
- Represents properties
of the thing to generate

Generator maps a random code to an output
- Generator is a (deep, convolutional) neural
network

GANs

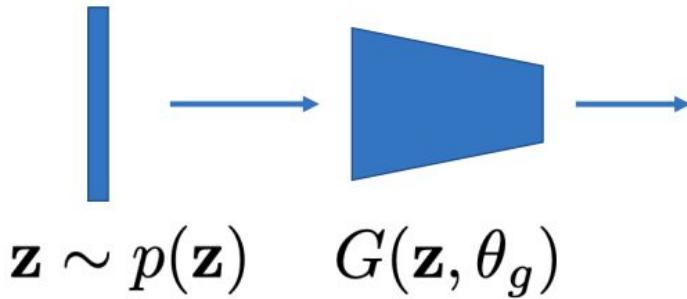


How do we learn the parameters of the generator?

No labeled data?

But lots of real unlabeled data...

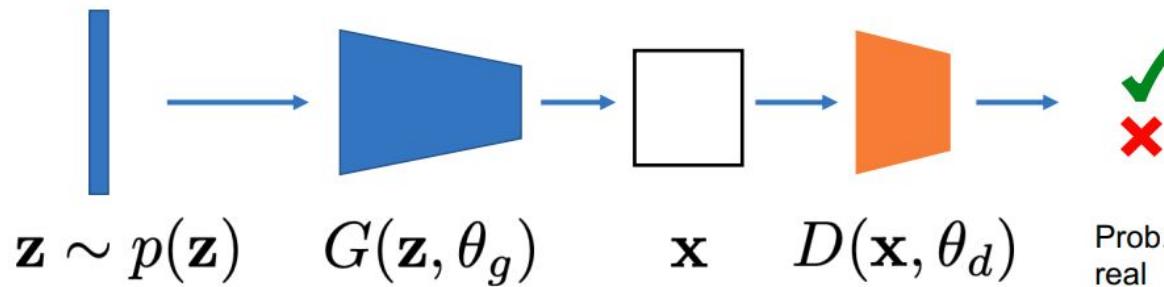
GANs



But which face should we generate?
What is a good output?
What if we generate non-faces?
Overfitting?
Variety?



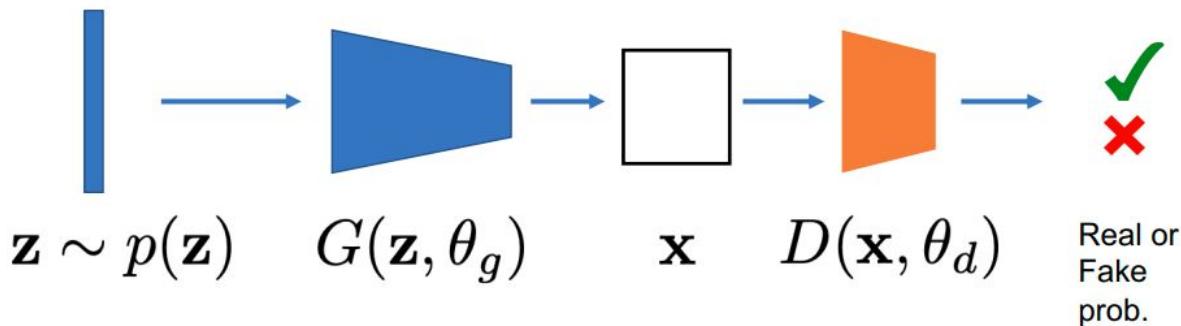
GANs



Given a discriminator, optimize the generator to fool the discriminator
- Generator should generate images that the discriminator thinks are real images, minimize wrt G:

$$\mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log(1 - D(G(\mathbf{z}, \theta_g), \theta_d))]$$

GANs

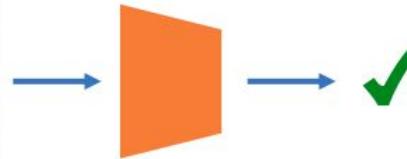


How do we train the discriminator?

GANs



$$\mathbf{x} \sim p_{data}(\mathbf{x})$$



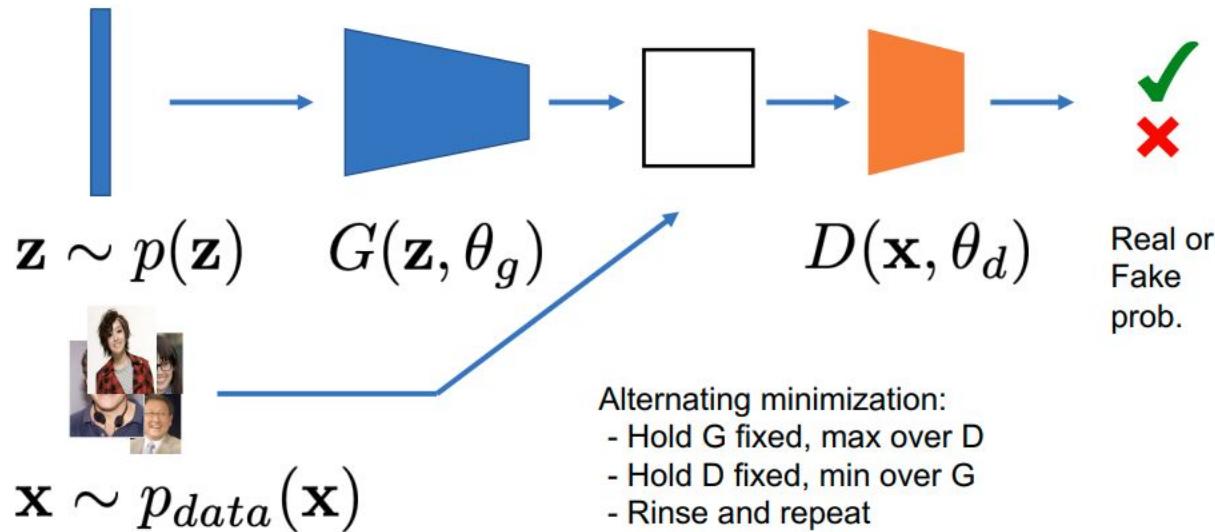
$$D(\mathbf{x}, \theta_d)$$

Should return real

Over the real images, make sure the discriminator thinks they are real
Maximize wrt. D:

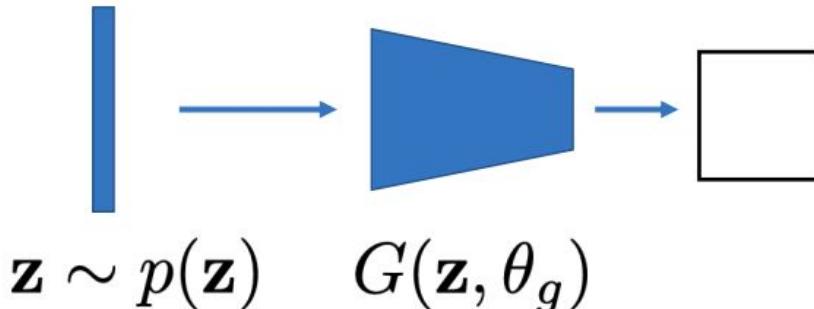
$$\mathbb{E}_{\mathbf{x} \sim p_{data}(\mathbf{x})} [\log D(\mathbf{x}, \theta_d)]$$

GANs



$$\min_{\theta_g} \max_{\theta_d} \mathbb{E}_{\mathbf{x} \sim p_{data}(\mathbf{x})} [\log D(\mathbf{x}, \theta_d)] + \mathbb{E}_{\mathbf{z} \sim p(\mathbf{z})} [\log(1 - D(G(\mathbf{z}, \theta_g), \theta_d))]$$

GANs



Once finished, discard the discriminator

Generator should generate realistic images for samples of \mathbf{z}

Alternating minimization:

- Hold G fixed, max over D
- Hold D fixed, min over G
- Rinse and repeat